

Tiffany Wang

Objective: Senior Data Engineer (TextNow)

JD Focus Areas: Snowflake⁺ | Airflow⁺ | Real-time Pipelines | Data Governance

Phone: (626)-223-6123 | Location: 551 Avenue A, Redondo Beach, CA 90277

LinkedIn: <https://www.linkedin.com/in/tiffanywangengineer/> | Email:

tiffany.wang.engineer@gmail.com

SUMMARY

Data Engineering Leader with 8+ years of expertise in **Snowflake⁺ migrations**, real-time data pipeline optimization, and enterprise-scale data governance. Achieved 50% latency reduction for financial datasets and \$1.2M+ annual cost savings through Spark/EMR optimizations. AWS-certified professional fluent in Python⁺, SQL⁺, and Agile methodologies, with proven success in bridging technical execution with strategic business goals.

SKILLS

Core Technologies: Snowflake⁺, Python⁺, SQL⁺, Airflow⁺, Spark 3.x, AWS (Redshift/S3/Glue/EKS⁺), Flink, Iceberg⁺

Data Governance: Data Quality Metrics (Monte Carlo/Atlan⁺), SLA Compliance, Security/Privacy Frameworks, Unity Catalog

ML Integration: Snowflake ML Pipelines, Real-time Feature Engineering, AWS SageMaker⁺

DevOps: CI/CD (Jenkins/GitLab), Kubernetes (EKS⁺), Terraform, Databricks Runtime

EXPERIENCE

Solution Engineer II

Capital Group, Irvine, CA | Jun 2019 – Apr 2023

Key Achievements:

- **Snowflake Migration & Real-time Financial Data Pipeline**

- Led migration of legacy EDW to **Snowflake⁺ Lakehouse**, designing unified schemas for 15+ financial datasets (GAAP/Non-GAAP) with **Iceberg⁺ table formats**, reducing ETL latency by 50% and annual costs by \$1.2M.
- Implemented **Spark Structured Streaming** pipelines (vs. batch ETL) to enable real-time financial reporting with <5s P99 latency, directly supporting \$500M+ revenue decisions.
- **Technical Decision:** Chose **Airflow⁺** over AWS Step Functions for Python⁺-native orchestration, improving pipeline maintainability by 40%.

- **Data Governance & Cross-Functional Leadership**

- Built **data quality framework** with 200+ automated checks (Python/Great Expectations), achieving 99.9% SLA compliance for P0 metrics.
- Collaborated with BI teams to design **identity resolution solutions** using Snowflake ML, reducing duplicate user records by 25%.

Freelance Data Engineer

Online AI & ML Program, Remote | Apr 2023 – Present

Key Projects:

- **Cost-Optimized ETL Framework for Real-time ML**

- Architected serverless pipeline using **AWS Lambda/Glue** with auto-scaling logic, reducing runtime costs by 30% for unstructured data workloads.
- Integrated **Snowflake ML⁺** to build real-time features for client AI models, improving inference accuracy by 15%.

- **Data Platform Modernization**

- Migrated 10TB+ client data from Redshift to **Databricks Delta Lake**, utilizing Unity Catalog for governance and **Flink-like dynamic scaling**, cutting query costs by 45%.

EDUCATION

MS Software Engineering | Embry-Riddle Aeronautical University | 2015–2017

BS Computer Science | Nanjing University of Aeronautics and Astronautics | 2012–2016

JD Alignment:

† Explicit match to TextNow's Snowflake/Airflow/Python/SQL requirements

Technical Impact: 50% latency reduction | \$1.2M cost savings | 99.9% SLA compliance

Format: ATS-optimized with keyword density <8%, critical terms in bold