

Tiffany Wang

Objective: Data Engineer - Zonic

Focus Areas: AWS EMR⁺ | Python (Pandas)⁺ | Real-time Data Pipelines | API-Driven Solutions

Phone: (626)-223-6123 | Location: Redondo Beach, CA 90277

LinkedIn: [linkedin.com/in/tiffanywangengineer](https://www.linkedin.com/in/tiffanywangengineer) | Email: tiffany.wang.engineer@gmail.com

SUMMARY

AWS-Certified Data Engineer with 8+ years of experience building **scalable data solutions** for financial and AI systems. Expert in Python⁺ (Pandas, PySpark), AWS⁺ (EMR, S3, Lambda), and SQL⁺, delivering 90% latency reduction and 99.99% uptime for mission-critical pipelines. Proven track record in GDPR-compliant data governance and cross-functional collaboration with 20+ engineering teams.

SKILLS

Core Stack:

Python⁺ (Pandas⁺, PySpark), AWS⁺ (EMR⁺, S3⁺, Lambda⁺, Athena), SQL⁺, Airflow⁺, Git⁺

Data Engineering: ETL Optimization, Data Modeling⁺, Entity Detection⁺, Record Linking⁺, Parquet/JSON/CSV

APIs & Governance: REST/SOAP APIs⁺, Snowflake⁺, Kafka, Dynamic Data Masking, OpenLineage

EXPERIENCE

Senior Cloud Data Engineer

Capital Group | Irvine, CA | 2019–2023

AWS EMR⁺-Kafka Real-Time Pipeline

- Built Spark Streaming system on AWS EMR⁺ processing **2M+ events/sec**, reducing NLP feature latency by **90%** (15min → 9s)
- **Technical Choice:** Selected Hudi over Parquet for delta streaming, improving time-travel query speed by **60%**
- Orchestrated pipelines via Airflow⁺-Temporal hybrid scheduler, achieving **99.99% uptime**

GDPR Compliance Platform (Snowflake⁺)

- Designed column-level encryption APIs with Snowflake⁺ Dynamic Data Masking, ensuring GDPR compliance in **3 months**
- **Hook:** Implemented automated data lineage tracking (OpenLineage), adopted by **20+ teams** for audit logging

AI Data Infrastructure Consultant

FinTech Innovation Lab | Remote | 2023–Present

Multi-Cloud Feature Store (AWS S3⁺ + Snowflake⁺)

- Architected hybrid feature store reducing serving latency by **75%** for a **50M+ user** fintech copilot
- **SDK Development:** Created Python⁺ SDK with Pandas⁺ integration, adopted by **150+ data scientists**

Anomaly Detection Mesh (Kafka-Debezium)

- Deployed CDC pipeline using Debezium→Kafka for real-time SLA monitoring, detecting **99.7% breaches**
- **Impact:** Reduced incident response time by **40%** via automated alerts

EDUCATION

MS Software Engineering | Carnegie Mellon University | 2015–2017

BS Computer Science | Tsinghua University | 2011–2015

Quantifiable Impact:

- 90% latency reduction | 99.99% uptime | \$500K+ cost savings from optimized cloud resources

† **JD-Aligned Keywords**

AWS EMR/S3/Lambda | Python/Pandas | Airflow | Snowflake | Entity Detection | GDPR Compliance